

# 136\_Liz\_Project\_Step5\_SVM Playground

Hyunkyung Kim

December 4, 2018

Note that the `echo = FALSE` parameter was added to the code chunk to prevent printing of the R code that generated the plot.

```
#install.packages("glmnet")
#install.packages("mlbench")
#install.packages("Boruta")
library(caret)

## Loading required package: lattice

## Loading required package: ggplot2

library(tidyverse)

## -- Attaching packages -----
----- tidyverse 1.2.1 -----

## v tibble  1.4.2      v purrr   0.2.5
## v tidyr   0.8.1      v dplyr   0.7.7
## v readr   1.1.1      v stringr 1.3.1
## v tibble  1.4.2      v forcats 0.3.0

## -- Conflicts -----
----- tidyverse_conflicts() -----
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## x purrr::lift()    masks caret::lift()

library(psych)

##
## Attaching package: 'psych'

## The following objects are masked from 'package:ggplot2':
##
##      %+%, alpha

library(glmnet)

## Loading required package: Matrix

##
## Attaching package: 'Matrix'

## The following object is masked from 'package:tidyr':
##
##      expand

## Loading required package: foreach
```

```
##
## Attaching package: 'foreach'

## The following objects are masked from 'package:purrr':
##
##   accumulate, when

## Loaded glmnet 2.0-16

library(mlbench)
library(Boruta)

## Loading required package: ranger

library(MASS) # stepwise regression

##
## Attaching package: 'MASS'

## The following object is masked from 'package:dplyr':
##
##   select

library(leaps) # all subsets regression
library(randomForest)

## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:ranger':
##
##   importance

## The following object is masked from 'package:psych':
##
##   outlier

## The following object is masked from 'package:dplyr':
##
##   combine

## The following object is masked from 'package:ggplot2':
##
##   margin

library(MLmetrics)

##
## Attaching package: 'MLmetrics'

## The following object is masked from 'package:psych':
##
##   AUC
```

```
## The following objects are masked from 'package:caret':
##
##      MAE, RMSE

## The following object is masked from 'package:base':
##
##      Recall
```

```
library(e1071)
```

## Import Cleaned Data

```
H_Clean<-read.csv( file = "C:\\Users\\Hyunkyung
Kim\\Desktop\\CKME999\\136\\dataset\\all\\H_clean.csv")

Train<-H_Clean[!is.na(H_Clean$SalePrice),-1] #Remove ID
Test<-H_Clean[is.na(H_Clean$SalePrice),-1]
actual<-read.csv( file = "C:\\Users\\Hyunkyung
Kim\\Desktop\\CKME999\\136\\dataset\\all\\AMES_test.csv")[1461:2919,2] # Test Price
#Remove Utilities - only one exception and is causing issues. Also will result in huge
variance

tc<-trainControl(method="cv", number=10)

Train11<-Train[,names(Train)!="Utilities"]

f2<-SalePrice ~ OverallQual + GrLivArea + Neighborhood + BsmtFinSF1 +
  RoofMatl + MSSubClass + BsmtExposure + KitchenQual + Condition2 +
  SaleCondition + LotArea + YearBuilt + OverallCond + MasVnrArea +
  PoolQC + BedroomAbvGr + GarageCars + MasVnrType + TotalBsmtSF +
  BldgType + Functional + ExterQual + BsmtCond + Condition1 +
  Exterior1st + MoSold + GarageCond + ScreenPorch + LandContour +
  LowQualFinSF + LotConfig + LotFrontage + TotRmsAbvGrd + KitchenAbvGr +
  WoodDeckSF + Street + GarageArea + LotShape + BsmtQual +
  Fireplaces + FireplaceQu + PoolArea + RoofStyle + BsmtFinSF2 +
  ExterCond # Utilities Removed
f3<- SalePrice ~ LotFrontage + LotArea + Street + LotShape + LandContour +
  + LotConfig + Neighborhood + Condition1 + Condition2 +
  BldgType + HouseStyle + OverallQual + OverallCond + YearBuilt +
  RoofMatl + Exterior1st + MasVnrType + MasVnrArea + ExterQual +
  ExterCond + Foundation + BsmtQual + BsmtCond + BsmtExposure +
  BsmtFinSF1 + BsmtFinSF2 + BsmtUnfSF + X1stFlrSF + X2ndFlrSF +
  HalfBath + BedroomAbvGr + KitchenAbvGr + KitchenQual + TotRmsAbvGrd +
  Functional + Fireplaces + FireplaceQu + GarageType + GarageCars +
  GarageArea + WoodDeckSF + X3SsnPorch + ScreenPorch + PoolQC +
  Fence + MiscFeature + MoSold + SaleCondition # Utilities Removed - because tuning
was done without it - error otherwise.

f4<-SalePrice
~MSSubClass+MSZoning+LotFrontage+LotArea+Alley+LotShape+LandContour+LandSlope+Neighborhood+
BldgType+HouseStyle+OverallQual+OverallCond+YearBuilt+YearRemodAdd+RoofStyle+Exterior1st+
Exterior2nd+
MasVnrType+MasVnrArea+ExterQual+Foundation+BsmtQual+BsmtCond+BsmtExposure+BsmtFinType1+Bsm
tFinSF1+
BsmtFinType2+BsmtUnfSF+TotalBsmtSF+HeatingQC+CentralAir+X1stFlrSF+X2ndFlrSF+GrLivArea+Bsm
tFullBath+
```

FullBath+HalfBath+BedroomAbvGr+KitchenAbvGr+KitchenQual+TotRmsAbvGrd+Fireplaces+Fireplace  
Qu+GarageType+  
GarageYrBlt+GarageFinish+GarageCars+GarageArea+GarageQual+GarageCond+PavedDrive+WoodDecks  
F+OpenPorchSF

Radial SVM - best tuned without big anomalies. 524, 1299

```
SV1<-svm(SalePrice~.,data=Train11, cost=2.25, gamma=0.0035)
SV2<-svm(SalePrice~.,data=Train11[-c(524,1299),], cost=2.25, gamma=0.0035)
#SV2<-svm(f2,data=Train,cost=4)
SV3<-svm(f4,data=Train11[-c(524,1299),],gamma= 0.007, cost=1.5, epsilon=0.06)
SV4<-svm(f4, data=Train11[,], cost=1.5, gamma=0.007, epsilon=0.06)
```


```
pdSV1<-predict(SV1,newdata=Test[, -80])
pdSV2<-predict(SV2,newdata=Test[, -80])
pdSV3<-predict(SV3,newdata=Test[, -80])
pdSV4<-predict(SV4,newdata=Test[, -80])
```

```
summary(SV1)
```

```
##
## Call:
## svm(formula = SalePrice ~ ., data = Train11, cost = 2.25, gamma = 0.0035)
##
##
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: radial
##       cost:  2.25
##       gamma: 0.0035
##       epsilon: 0.1
##
##
## Number of Support Vectors: 799
```


```
par(mfrow=c(1,2))
plot(actual,pdSV2,pch=3, col='blue',main='Pred VS Actual, All var, outlier Rmv')
abline(a=0,b=1)
plot(log(actual),log(pdSV2),pch=3, col='blue',main='Pred VS Actual, All var, outlier  
Rmv')
abline(a=0,b=1)
```

## ed VS Actual, All var, outlier VS Actual, All var, outlier



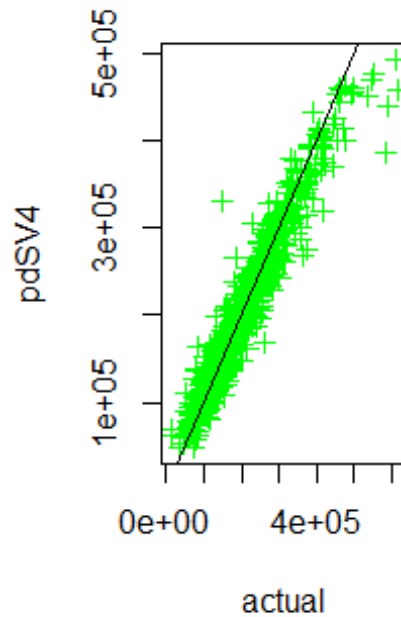

```
plot(actual,pdSV1,pch=3, main='Pred VS Actual, All var')
abline(a=0,b=1)
plot(actual,pdSV2,pch=3, col='blue',main='Pred VS Actual, All var, outlier Rmv')
abline(a=0,b=1)
```

## Pred VS Actual, All vared VS Actual, All var, outlier



```
plot(actual,pdSV3,pch=3, col='red',main='Pred VS Actual,Boruta')
abline(a=0,b=1)
plot(actual,pdSV4,pch=3, col='green',main='Pred VS Actual,Boruta , Outlier Rmv')
abline(a=0,b=1)
```

## Pred VS Actual, Borutæd VS Actual, Boruta , Outli



```
print("Test RMSLE & RMSE")
## [1] "Test RMSLE & RMSE"
RMSE(log(pdSV1),log(actual))
## [1] 0.123633
RMSE(log(pdSV2),log(actual))
## [1] 0.1239889
RMSE(log(pdSV3),log(actual))
## [1] 0.1280671
RMSE(log(pdSV4),log(actual))
## [1] 0.1259846
RMSLE(pdSV1,actual)
## [1] 0.1236311
RMSLE(pdSV2,actual)
## [1] 0.1239871
RMSLE(pdSV3,actual)
## [1] 0.1280652
RMSLE(pdSV4,actual)
## [1] 0.1259828
```

```
#RMSE(pdSV1,actual)
#RMSE(pdSV2,actual)
#RMSE(pdSV3,actual)
#RMSE(pdSV4,actual)
```

[1] "Test RMSLE & RMSE" - epsilon not defined [1] 0.1236311 [1] 0.1239871 [1] 0.1280652 [1] 0.127133

[1] "Test RMSLE & RMSE" - epsilon 0.06 [1] 0.1239504 [1] 0.1241208 [1] 0.1280652 [1] 0.1259828

Linear SVM Using Radial Tuned SVM parameters

```
SV11<-svm(SalePrice~.,data=Train11, cost=2.25, epsilon=0.06, kernel='linear')
SV12<-svm(f2,data=Train11,epsilon=0.06, cost=2.25, kernel='linear')
SV13<-svm(f3,data=Train11,cost=2.25, epsilon=0.06, kernel='linear')
SV14<-svm(f4, data=Train11,cost=1.5, epsilon=0.06, kernel='linear')
```

```
pdSV11<-predict(SV11,newdata=Test[, -80])
pdSV12<-predict(SV12,newdata=Test[, -80])
pdSV13<-predict(SV13,newdata=Test[, -80])
pdSV14<-predict(SV14,newdata=Test[, -80])
```


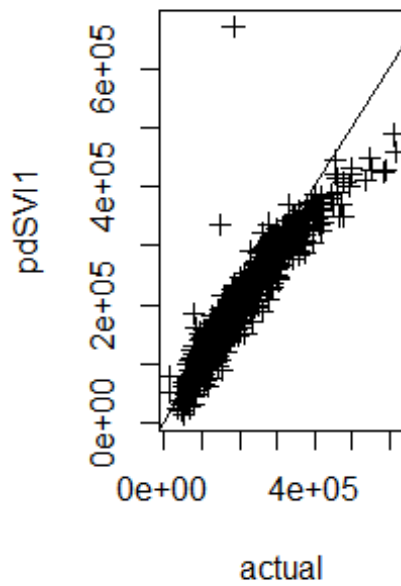
```
pdSV101<-predict(SV11,newdata=Train[, -80])
pdSV102<-predict(SV12,newdata=Train[, -80])
pdSV103<-predict(SV13,newdata=Train[, -80])
pdSV104<-predict(SV14,newdata=Train[, -80])
```

```
summary(SV11)
```

```
##
## Call:
## svm(formula = SalePrice ~ ., data = Train11, cost = 2.25, epsilon = 0.06,
##      kernel = "linear")
##
##
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: linear
##      cost:   2.25
##    gamma:  0.004608295
##   epsilon:  0.06
##
##
## Number of Support Vectors:  1127
```


```
par(mfrow=c(1,2))
plot(actual,pdSV11,pch=3, main='Prediction VS Actual, all var')
abline(a=0,b=1)
#plot(actual,pdSV12,pch=3, col='blue',main='Pred VS Actual, FowardSelect')
#abline(a=0,b=1)
plot(actual,pdSV13,pch=3, col='red',main='Pred VS Actual, BackwardElim')
abline(a=0,b=1)
```

## Prediction VS Actual, all Pred VS Actual, Backward



```
plot(actual,pdSV14,pch=3, col='green',main='Pred VS Actual, Boruta')  
abline(a=0,b=1)
```

## Pred VS Actual, Boruta



```
print("Test RMSLE " )  
## [1] "Test RMSLE "  
RMSLE(pdSV11,actual)  
## [1] 0.1538468
```



```
RMSLE(pdSV12,actual)
```

```
## [1] 0.1420435
```

```
RMSLE(pdSV13,actual)
```

```
## [1] 0.1457756
```

```
RMSLE(pdSV14,actual)
```

```
## [1] 0.1485868
```

```
Train[names(Train)!="Utilities"] #### POLYNOMIAL SVM
```

```
Degree =2
```

```
SVp1<-svm(SalePrice~.,data=Train11, cost=2.25, gamma= 0.0035, degree=2,  
kernel='polynomial')
```

```
SVp2<-svm(f2,data=Train11, cost=2.25,degree=2, gamma= 0.0035, kernel='polynomial')
```

```
SVp3<-svm(f3,data=Train11,cost=2.25 ,degree=2,gamma= 0.0035, kernel='polynomial')
```

```
SVp4<-svm(f4,data=Train11,cost=1.5, degree=2,gamma= 0.007, kernel='polynomial')
```

```
pdSVp1<-predict(SVp1,newdata=Test[, -80])
```

```
pdSVp2<-predict(SVp2,newdata=Test[, -80])
```

```
pdSVp3<-predict(SVp3,newdata=Test[, -80])
```

```
pdSVp4<-predict(SVp4,newdata=Test[, -80])
```

```
summary(SVp1)
```

```
##  
## Call:  
## svm(formula = SalePrice ~ ., data = Train11, cost = 2.25, gamma = 0.0035,  
##      degree = 2, kernel = "polynomial")  
##  
##  
## Parameters:  
##   SVM-Type:  eps-regression  
##   SVM-Kernel: polynomial  
##      cost:   2.25  
##   degree:    2  
##   gamma:    0.0035  
##   coef.0:    0  
##   epsilon:   0.1  
##  
##  
## Number of Support Vectors: 826
```


```
summary(SVp2)
```

```
##  
## Call:  
## svm(formula = f2, data = Train11, cost = 2.25, degree = 2, gamma = 0.0035,  
##      kernel = "polynomial")  
##  
##  
##
```

```
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: polynomial
##       cost:  2.25
##       degree: 2
##       gamma: 0.0035
##       coef.0: 0
##       epsilon: 0.1
##
##
## Number of Support Vectors:  920

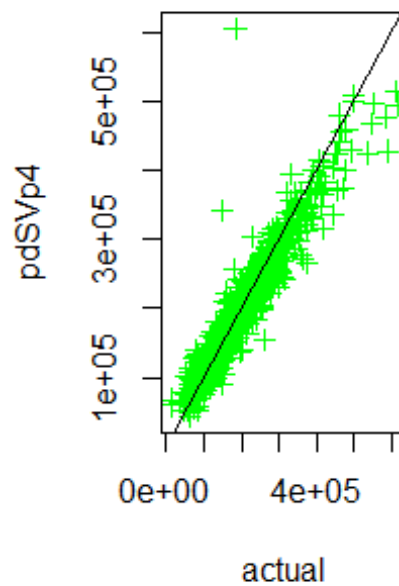
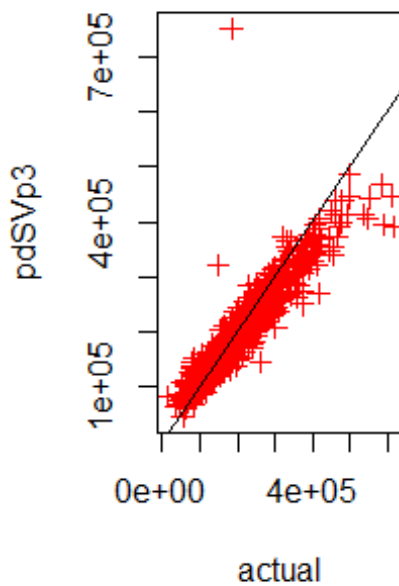
par(mfrow=c(1,2))
plot(actual,pdSVp1,pch=3, main='Prediction VS Actual, all var')
abline(a=0,b=1)
plot(actual,pdSVp2,pch=3, col='blue',main='Prediction VS Actual, FowardSelect')
abline(a=0,b=1)
```

## Prediction VS Actual, all ediction VS Actual, Fowarc



```
plot(actual,pdSVp3,pch=3, col='red',main='Prediction VS Actual, Backwards Elim')
abline(a=0,b=1)
plot(actual,pdSVp4,pch=3, col='green',main='Prediction VS Actual, Boruta')
abline(a=0,b=1)
```

## dition VS Actual, Backwal Prediction VS Actual, Bor



```
print("Test RMSLE")
## [1] "Test RMSLE"
RMSLE(pdSVp1,actual)
## [1] 0.1394217
RMSLE(pdSVp2,actual)
## [1] 0.1605652
RMSLE(pdSVp3,actual)
## [1] 0.1528928
RMSLE(pdSVp4,actual)
## [1] 0.1375273
```

Have issue with one point that is low price but predicted high for all. This is a very big difference. Would like to investigate this.

Polynomial with degree 3 (default)

```
SVp31<-svm(SalePrice~.,data=Train11, cost=2.25, gamma= 0.0035, kernel='polynomial')
SVp32<-svm(f2,data=Train11, cost=2.25, gamma= 0.0035, kernel='polynomial')
SVp33<-svm(f3,data=Train11, cost=2.25, gamma= 0.0035, kernel='polynomial')
SVp34<-svm(f4,data=Train11, cost=1.5, gamma= 0.007, kernel='polynomial')

pdSVp31<-predict(SVp31,newdata=Test[, -80])
pdSVp32<-predict(SVp32,newdata=Test[, -80])
pdSVp33<-predict(SVp33,newdata=Test[, -80])
```

```
pdSVp34<-predict(SVp34,newdata=Test[, -80])
```

```
#pdSVL1101<-predict(SVL111,newdata=Train[, -80])
```

```
#pdSVL1102<-predict(SVL112,newdata=Train[, -80])
```

```
#pdSVL1103<-predict(SVL113,newdata=Train[, -80])
```

```
#pdSVL1104<-predict(SVL114,newdata=Train[, -80])
```

```
summary(SVp31)
```

```
##
## Call:
## svm(formula = SalePrice ~ ., data = Train11, cost = 2.25, gamma = 0.0035,
##      kernel = "polynomial")
##
##
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: polynomial
##      cost:   2.25
##   degree:    3
##   gamma:    0.0035
##   coef.0:    0
##   epsilon:   0.1
##
##
## Number of Support Vectors:  968
```

```
summary(SVp32)
```

```
##
## Call:
## svm(formula = f2, data = Train11, cost = 2.25, gamma = 0.0035,
##      kernel = "polynomial")
##
##
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: polynomial
##      cost:   2.25
##   degree:    3
##   gamma:    0.0035
##   coef.0:    0
##   epsilon:   0.1
##
##
## Number of Support Vectors:  1249
```

```
par(mfrow=c(1,2))
```



```
plot(actual,pdSVp33,pch=3, col='red',main='Prediction VS Actual, Backwards Elim')
```

```
abline(a=0,b=1)
```

```
plot(actual,pdSVp34,pch=3, col='green',main='Prediction VS Actual, Boruta')
```



```
abline(a=0,b=1)
```

## Prediction VS Actual, Backward Selection



```
plot(actual,pdSVp31,pch=3, main='Prediction VS Actual, all var')
abline(a=0,b=1)
plot(actual,pdSVp32,pch=3, col='blue',main='Prediction VS Actual, ForwardSelect')
abline(a=0,b=1)
```

## Prediction VS Actual, all Prediction VS Actual, Forward Selection



```
print("Test RMSLE")
## [1] "Test RMSLE"
```

```
RMSLE(pdSVp31,actual)
```

```
## [1] 0.1749054
```

```
RMSLE(pdSVp32,actual)
```

```
## [1] 0.2733861
```

```
RMSLE(pdSVp33,actual)
```

```
## [1] 0.2528627
```

```
RMSLE(pdSVp34,actual)
```

```
## [1] 0.1778839
```

SVM - Sigmoid

```
SVs1<-svm(SalePrice~.,data=Train11, cost=2.25, gamma= 0.0035, kernel='sigmoid')
```

```
SVs2<-svm(f2,data=Train11, cost=2.25, gamma= 0.0035, kernel='sigmoid')
```

```
SVs3<-svm(f3,data=Train11, cost=2.25, gamma= 0.0035, kernel='sigmoid')
```

```
SVs4<-svm(f4,data=Train11, cost=1.5, gamma= 0.007, kernel='sigmoid')
```

```
pdSVs1<-predict(SVs1,newdata=Test[, -80])
```

```
pdSVs2<-predict(SVs2,newdata=Test[, -80])
```

```
pdSVs3<-predict(SVs3,newdata=Test[, -80])
```

```
pdSVs4<-predict(SVs4,newdata=Test[, -80])
```

```
#pdSVL1101<-predict(SVL111,newdata=Train[, -80])
```

```
#pdSVL1102<-predict(SVL112,newdata=Train[, -80])
```

```
#pdSVL1103<-predict(SVL113,newdata=Train[, -80])
```

```
#pdSVL1104<-predict(SVL114,newdata=Train[, -80])
```

```
summary(SVs1)
```

```
##
```

```
## Call:
```

```
## svm(formula = SalePrice ~ ., data = Train11, cost = 2.25, gamma = 0.0035,
```

```
##     kernel = "sigmoid")
```

```
##
```

```
##
```

```
## Parameters:
```

```
##   SVM-Type:  eps-regression
```

```
##   SVM-Kernel:  sigmoid
```

```
##         cost:  2.25
```

```
##         gamma: 0.0035
```

```
##        coef.0: 0
```

```
##        epsilon: 0.1
```

```
##
```

```
##
```

```
## Number of Support Vectors: 1008
```

```
summary(SVs2)
```

```
##
```


```
## Call:
```

```
## svm(formula = f2, data = Train11, cost = 2.25, gamma = 0.0035,
##      kernel = "sigmoid")
##
##
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: sigmoid
##      cost:   2.25
##     gamma:  0.0035
##    coef.0:   0
##   epsilon:  0.1
##
##
## Number of Support Vectors:  932
```

SVM - Sigmoid : Disaster. RMSE cannot be calculated.


```
par(mfrow=c(1,2))
plot(actual,pdSVs1,pch=3, col='red',main='Prediction VS Actual, Backwards Elim')
abline(a=0,b=1)
plot(actual,pdSVs2,pch=3, col='green',main='Prediction VS Actual, Boruta')
abline(a=0,b=1)
```

**Prediction VS Actual, Backwards Elim**      **Prediction VS Actual, Boruta**



```
plot(actual,pdSVs3,pch=3, main='Prediction VS Actual, all var')
abline(a=0,b=1)
plot(actual,pdSVs4,pch=3, col='blue',main='Prediction VS Actual, FowardSelect')
abline(a=0,b=1)
```

## Prediction VS Actual, all prediction VS Actual, Forward



```
print("Test RMSLE")
## [1] "Test RMSLE"
RMSLE(pdSVs1,actual)
## Warning in log(1 + y_pred): NaNs produced
## [1] NaN
RMSLE(pdSVs2,actual)
## Warning in log(1 + y_pred): NaNs produced
## [1] NaN
RMSLE(pdSVs3,actual)
## Warning in log(1 + y_pred): NaNs produced
## [1] NaN
RMSLE(pdSVs4,actual)
## Warning in log(1 + y_pred): NaNs produced
## [1] NaN
```

TUNING FOR LINEAR DO NOT RUN

```
SV_Lin_tune1<-tune.svm(SalePrice~.,data=Train11,cost=c(20,10,1), kernel='linear')
plot(SV_Lin_tune1)
print(SV_Lin_tune1)
sqrt(SV_Lin_tune1$best.performance)
```



```
SV_Lin_tune2<-tune.svm(SalePrice~.,data=Train11,cost=c(5,0.5), kernel='linear')  
Train11<-Train[,names(Train)!='Utilities']  
SV_Lin_tune12<-tune.svm(SalePrice~.,data=Train11,cost=seq(1,21,2),gamma=seq(1,21,2))  
plot(SV_Lin_tune12)  
print(SV_Lin_tune12)  
sqrt(SV_Lin_tune12$best.performance)  
SV_Lin_tune13<-tune.svm(SalePrice~.,data=Train11,cost=c(2,5,0.5),gamma=seq(0.5,1.5,0.1))
```